Infosys Java Developer Interview Question(3-8 Yrs Experience)

Q1. How do you perform a filter + map + reduce in one stream pipeline?

```
Ans- al.stream().filter(I - > I %2==0).map(I-> I * I) .reduce(0,Integer::sum);
```

Q2. What is the difference between reduce() and collect() in java streams?

Ans - 1. Purpose

- reduce() → Used for mathematical-style reduction: sum, min, max, product, concatenation, etc. Produces a single value.
- collect() → Used for mutable reductions: collecting into collections (List, Set, Map) or building complex results. Produces a mutable container.

2. Return Type

- reduce() → returns a single value (or an Optional<T> if no identity is given).
- collect() → returns a container
 or aggregate (e.g., List, Set, Map, or even a string).

Q3. What is the difference between Collectors.groupingBy() and Collectors.PartioningBy()?

Ans. 1. Collectors.groupingBy()

- **Definition**: Groups elements of the stream by a *classifier function* (any mapping function).
- Return type: Map<K, List<T>> (or Map<K, Set<T>>, etc. if you give a downstream collector).
- Use case: When you want to group data into multiple categories.
- Example:

```
List<String> words =
List.of("apple", "banana",
"cherry", "apricot", "blueberry");
```

```
Map<Character, List<String>>
groupedByFirstLetter =
    words.stream()
        .collect(Collectors.group
ingBy(word -> word.charAt(0)));

System.out.println(groupedByFirstLetter);
// {a=[apple, apricot], b=[banana, blueberry], c=[cherry]}
Here, elements are grouped into many buckets, depending on the key.
```

2. Collectors.partitioningBy()

- **Definition**: A *special case of groupingBy()* where the classifier is a **predicate** (boolean test).
- Return type: Map<Boolean, List<T>>.
- Use case: When you want to divide data into two groups (true/false).
- Example:

```
List<String> words =
List.of("apple", "banana",
"cherry", "apricot", "blueberry");

Map<Boolean, List<String>>
partitionedByLength =
    words.stream()
        .collect(Collectors.partitioningBy(word -> word.length() >
5));

System.out.println(partitionedByLength);
// {false=[apple], true=[banana, cherry, apricot, blueberry]}
Here, elements are split into exactly two buckets:
```

- true → words with length > 5
- false \rightarrow words with length ≤ 5

Q4. How you would convert a List<Employee> into a Map<Department,List<Empl

oyee>> using streams?

Ans.

Map<Department,List<Employee>> mapResult = al.stream().collect(Collector s.groupingBy(Employee::ge t dept));

Q5.What is the difference between run() and start()? Ans.

1. start() method

- Declared in Thread class.
- Creates a **new thread of execution**.
- Calls the run () method in a separate call stack.
- Once you call start(), the JVM schedules the thread for execution.

```
class MyThread extends Thread {
    public void run() {
        System.out.println("Running in:
+ Thread.currentThread().getName());
}
public class Test {
    public static void main(String[]
args) {
        MyThread t1 = new MyThread();
        tl.start(); // starts a new
thread
        System.out.println("Main thread:
 + Thread.currentThread().getName());
    }
Possible Output:
Main thread: main
Running in: Thread-0
The Notice: run() is executed in a different thread
(Thread-0).
```

2. run() method

- Declared in Runnable interface (implemented/overridden in Thread subclass).
- Defines the code that will run when the thread starts.

• If you call run() directly, it is treated like a **normal method call**, **not** a separate thread.

Example:

```
MyThread t2 = new MyThread();
t2.run(); // just a normal method call
System.out.println("Main thread: " +
Thread.currentThread().getName());
Output:
```

```
Running in: main
Main thread: main
Here, run() executed in the main thread, not a new one.
```

Q6.What is synchronisation? Why is it needed?

Ans

In Java, synchronization is the process of controlling access to shared resources (like variables, objects, files) by multiple threads.

It ensures that **only one thread can access a shared resource at a time**, preventing inconsistent or corrupted data.

In simple terms:

Synchronization = "one thread at a time" rule for shared resources.

. Why is Synchronization Needed?

Because Java supports multithreading:

- Multiple threads may run **concurrently**.
- If threads modify the same shared resource **without control**, problems occur (called *race conditions*).

Example of Problem (Without Synchronization)

```
class Counter {
    private int count = 0;
    public void increment() {
        count++; // not atomic → multiple
threads may interfere
    public int getCount() { return count; }
}
public class Test {
    public static void main(String[] args)
throws InterruptedException {
        Counter counter = new Counter();
        Thread t1 = new Thread(() -> {
            for (int i = 0; i < 1000; i++)
counter.increment();
        });
        Thread t2 = new Thread(() \rightarrow {
            for (int i = 0; i < 1000; i++)
counter.increment();
```

```
    t1.start();
    t2.start();
    t1.join();
    t2.join();

    System.out.println("Final count: " +
counter.getCount());
    }
}
Expected Output:

2000
But Actual Output may be:

1984, 1992, etc. (varies)
    Because both threads update count at the same time, overwriting each other.
```

3. Fix Using Synchronization

```
class Counter {
    private int count = 0;

    public synchronized void increment() {
        count++; // synchronized ensures one
thread at a time
    }

    public int getCount() { return count; }
}
Now the result will always be:
```

Final count: 2000

Key Points

- Synchronization ensures thread safety.
- Prevents **race conditions** and data inconsistency.
- Achieved in Java via:
 - synchronized keyword (methods/blocks).
 - Locks (ReentrantLock, ReadWriteLock).
 - Concurrent utilities (ConcurrentHashMap, AtomicInteger).

Think of synchronization like a **traffic signal at a busy intersection**:

- Without it \rightarrow chaos (accidents).
- With it \rightarrow safe, orderly access, even if slower.

Q7. What is false sharing? Ans.

False sharing happens when multiple threads modify different variables that happen to lie close together in memory (same CPU cache line).

Even though the variables are *independent*, the CPU cache system treats them as one block

(cache line).

This causes unnecessary cache invalidations and slows down performance.

Why It Happens

- Modern CPUs don't load variables one by one.
- They load data in chunks called cache lines (usually 64 bytes).
- If two threads work on different variables that fall into the same cache line:
 - One thread's update invalidates the cache line for the other.
 - CPU keeps "bouncing" cache lines between cores.
 - This creates contention even though there's no real data dependency.

Illustration

Example (False Sharing Scenario):

```
class Data {
    volatile long x; //
Thread 1 writes here
    volatile long y; //
Thread 2 writes here
}
Data data = new Data();
// Thread 1
for (int i = 0; i < 1_000_000;
i++) {
    data.x++;
}
// Thread 2
for (int i = 0; i < 1 000 000;
i++) {
    data.y++;
}
 • x and y are different variables.
```

• But they might be placed in the same

cache line.

• So both threads unintentionally slow each other down.

How to Avoid False Sharing

1 Padding: Add unused fields to separate hot variables into different cache lines.

```
class Data {

volatile long x;

// padding to push y
into a different cache line
long p1, p2, p3, p4,
p5, p6, p7;
volatile long y;

}
```

8 2)@Contended annotation (Java 8+, but

```
needs JVM option -XX:-
RestrictContended):
```

```
import
 jdk.internal.vm.annotation.
 Contended;
9
10
    class Data {
11
        @Contended
12
        volatile long x;
13
14
        @Contended
15
        volatile long y;
16
    }
17
```

Tip: Use concurrency-friendly data structures like LongAdder or ConcurrentHashMap that internally avoid false sharing.

Key Takeaways

- False sharing = performance bug, not correctness bug.
- Happens when independent variables share the same cache line.
- Causes unnecessary cache invalidation between threads.
- Fixed by **padding**, @Contended, or better concurrency utilities.

Think of it like two people trying to write on two pages of the same notebook:

• They're not editing the same sentence, but the notebook keeps being taken away and passed back, slowing both down.

Q8. How would you implement a producer-consumer problem?

Ans.

Producer → generates data and puts it in a shared buffer.

Consumer \rightarrow takes data from the buffer and processes it.

The buffer has limited capacity.

We must ensure:

- Producers don't add to a *full* buffer.
- Consumers don't take from an *empty buffer*.
- No race conditions (thread safety).

Q9.What is the difference

between synchronized and Reentrant lock? Ans

1. synchronized (Keyword)

- Language-level construct (built into JVM).
- Used to lock a **method** or a **block**.
- Implicitly acquires and releases the lock.
- Once the thread exits the block/method, the lock is automatically released (even if an exception occurs).
- Always **reentrant** (a thread can acquire the same lock again without deadlock).

```
public synchronized void increment()
{
    count++;
}

public void decrement() {
    synchronized (this) {
       count--;
    }
}
```

2. ReentrantLock (Class in java.util.concurrent.locks)

- **Explicit lock class** introduced in Java 5.
- Provides greater flexibility than synchronized.
- Locking/unlocking must be done **manually** with lock() and unlock().
- Supports **fairness policy** (queue threads FIFO).
- Offers **tryLock()** (non-blocking attempt) and **lockInterruptibly()** (responds to interrupts).
- Must call unlock() in a finally block to avoid deadlocks.

```
import
java.util.concurrent.locks.ReentrantL
ock;

ReentrantLock lock = new
ReentrantLock();

public void increment() {
    lock.lock();
    try {
        count++;
    } finally {
```

```
lock.unlock(); // must
release manually
}
```

When to Use What?

- Use synchronized:
 - When synchronization needs are simple.
 - For most common thread-safe methods/blocks.
 - Cleaner and less error-prone.
- Use ReentrantLock:
 - When you need fairness, tryLock(), or interruptible locks.
 - When you need multiple condition variables for more complex producer—consumer logic.
 - When fine-grained control is required.

Think of it like this:

- synchronized = automatic car (easy, but less flexible).
- ReentrantLock = manual car \rightleftharpoons (more control,

Q10. What is the volatile keyword? Ans

- volatile is a **Java keyword** used with **variables**.
- It tells the JVM and CPU that the variable's value may be modified by multiple threads.
- Ensures that **reads/writes to that variable are directly done in main memory**, not cached in CPU registers or thread-local caches.

In simple words: volatile guarantees visibility of changes to variables across threads.

2. Key Points About volatile

1 Visibility Guarantee

- If Thread A updates a **volatile** variable, Thread B will see the updated value immediately.
- Without volatile, Thread B may see a stale cached value.

2 Atomicity

- volatile does NOT guarantee atomicity for compound operations like count++.
- Only guarantees visibility of single reads/writes.

3 When to Use

- For **flags or state variables** shared between threads.
- Example: boolean running = true; used to stop a thread.

Q11. What is N+1 select problem?

Ans

The N+1 Select Problem is a common performance issue in ORM frameworks like Hibernate or JPA when fetching related entities from a database. It can cause a large number of SQL queries and drastically slow down your application.

Let me break it down step by step:

1. The Problem

Suppose you have two entities: Author and

Book, where an Author has many Books.

```
@Entity
class Author {
    @Id
    private Long id;
    private String name;

    @OneToMany(mappedBy = "author")
    private List<Book> books;
}

@Entity
class Book {
    @Id
    private Long id;
    private String title;

    @ManyToOne
    private Author author;
}
```

Now, you want to fetch all authors **and their books**:

```
List<Author> authors = authorRepository.findAll();
for (Author author: authors) {
    System.out.println(author.getBooks().size());
}
```

2. What Actually Happens

1 The ORM executes **1 query** to fetch all authors:

```
SELECT * FROM author;
```

2

3 Then, for each author, it executes 1 query to fetch their books:

```
SELECT * FROM book WHERE author_id = 1;

4 SELECT * FROM book WHERE author_id = 2;

5 SELECT * FROM book WHERE author_id = 3;

6 ...

7
```

If you have **N** authors, you run **1** + **N** queries. Hence the name **N+1** select problem.

3. Why It's Bad

• If N is large, it can cause **hundreds or thousands of queries**, drastically slowing down the application.

Wastes database resources and increases latency.

4. How to Solve It

Solution 1: Eager Fetch / Join Fetch Use Join Fetch to load authors and books in one query:

```
@Query("SELECT a FROM Author a JOIN FETCH
a.books")
List<Author> findAllAuthorsWithBooks();
Generates a single SQL query:
```

```
SELECT a.*, b.*
FROM author a
JOIN book b ON a.id = b.author_id;
```

Solution 2: Batch Fetching

Configure the ORM to fetch collections in batches instead of one by one:

```
@BatchSize(size = 10)
private List<Book> books;
```

Key Takeaway: The N+1 select problem happens when your ORM executes 1 query for the main entity + N queries for each related entity, which can be avoided with join fetches, batch fetching, or projections.

Q12. How do you implement inheritance in JPA?

Ans

In **JPA**, inheritance allows you to map an **object-oriented class hierarchy** to database tables. JPA provides three main strategies for inheritance mapping. Let's go step by step.

1. Single Table Inheritance (SINGLE_TABLE)

- All classes in the hierarchy are stored in a single table.
- A discriminator column distinguishes the type of each row.

```
@Entity
@Inheritance(strategy =
InheritanceType.SINGLE TABLE)
```

```
@DiscriminatorColumn(name = "vehicle_type")
public abstract class Vehicle {
    @Id
    private Long id;
    private String manufacturer;
}
@Entity
public class Car extends Vehicle {
    private int seatingCapacity;
}
@Entity
public class Bike extends Vehicle {
    private boolean hasCarrier;
Generated Table: Vehicle
i
d
Pros: Simple, single table, good performance
```

2. Joined Table Inheritance (JOINED)

• Each class has its **own table**.

Cons: Many nullable columns for subclasses

• The **child table references the parent table** via primary key.

```
@Entity
@Inheritance(strategy =
InheritanceType.JOINED)
public abstract class Vehicle {
```

```
@Id
    private Long id;
    private String manufacturer;
}

@Entity
public class Car extends Vehicle {
    private int seatingCapacity;
}

@Entity
public class Bike extends Vehicle {
    private boolean hasCarrier;
}
```

Generated Tables:

- Vehicle → id, manufacturer
- Car → id (FK to Vehicle), seatingCapacity
- Bike → id (FK to Vehicle), hasCarrier

Pros: Normalized, no null columns

Cons: Joins needed to fetch child entities (slower)

3. Table Per Class (TABLE PER CLASS)

- Each concrete class has its **own table** with all fields of parent included.
- No parent table is used for storing data.

```
@Entity
@Inheritance(strategy =
InheritanceType.TABLE_PER_CLASS)
```

```
public abstract class Vehicle {
    @Id
    private Long id;
    private String manufacturer;
}

@Entity
public class Car extends Vehicle {
    private int seatingCapacity;
}

@Entity
public class Bike extends Vehicle {
    private boolean hasCarrier;
}
```

Generated Tables:

- Car → id, manufacturer, seatingCapacity
- Bike → id, manufacturer, hasCarrier

Pros: Simple, no joins

Cons: Queries across hierarchy are complex; IDs must be unique

across all tables

Key Annotations

- @Inheritance(strategy = ...) → defines the inheritance strategy
- @DiscriminatorColumn → used in SINGLE_TABLE for type differentiation
- @DiscriminatorValue("Car") → optional, specify value for child type

▼ Tip:

- **SINGLE_TABLE** → best for performance if subclasses have few additional fields
- **JOINED** → best for normalized schema
- TABLE_PER_CLASS → rarely used, mainly for abstract parent with very different subclasses

Q13. How do you write a native SQL query in JPA?

Ans

1. Using @Query with nativeQuery = true
When using Spring Data JPA:

```
public interface
UserRepository extends
JpaRepository<User, Long> {
```

```
= :email",
        nativeQuery = true
    User
findByEmail(@Param("email")
String email);
}
Here:
```

- - value = the SQL query
 - nativeQuery = true tells JPA it's raw SQL, not JPQL

2. Using @NamedNativeQuery (on Entity)

You can define a named native query on your entity class:

```
@Entity
@NamedNativeQuery(
    name =
"User.findByStatus",
    query = "SELECT * FROM
```

```
users WHERE status = ?1",
    resultClass = User.class
public class User {
    @Id
    private Long id;
    private String email;
    private String status;
}
Then call it like:
List<User> users =
entityManager
         .createNamedQuery("Use
r.findByStatus", User.class)
         .setParameter(1,
"ACTIVE")
         .getResultList();
3. Using
EntityManager.createNativeQuery()
If you want more flexibility (manual use):
```

```
List<Object[]> results =
entityManager
    .createNativeQuery("SELECT
u.id, u.email FROM users u
WHERE u.status = ?")
    .setParameter(1, "ACTIVE")
    .getResultList();

for (Object[] row : results) {
    Long id = ((Number))
row[0]).longValue();
    String email = (String)
row[1];
}
```